

Consensus Number

魏恒峰

hfwei@nju.edu.cn

2017 年 12 月 14 日




Feel Free to
Ask Questions



Are you Familiar with **Concurrent Programming**?



The Key: **Synchronization!**



Using the **Synchronization Primitives**
Provided by Your Favorite Languages.

synchronized
Semaphore

BlockingQueue
ConcurrentMap

Phaser
Barrier

synchronized
Semaphore

BlockingQueue
ConcurrentMap

Phaser
Barrier

```
class AtomicInteger:
    get()
    set(int newValue)

    getAndIncrement()
    getAndDecrement()
    getAndSet(int newValue)

    compareAndSet(int expectedValue, int newValue)
```

```
compareAndSet(int expectedValue, int newValue)  
compareAndSwap(int expectedValue, int newValue)
```

CAS — CMPXCHG

`impl.`

`usage`

Tasks: Implementing Consensus using Given Primitives.

MISSION:
POSSIBLE

MISSION:
IMPOSSIBLE



Consensus



"It looks like we have a consensus."

"It looks like we have a consensus."

Propose



Decide

Propose



Decide

Definition (The Consensus Problem)

Agreement All (non-faulty) processes must **agree on the same value**.

Validity The common decision value must be the value **proposed** by some process.

Termination Each (non-faulty) process must **eventually** decide on a value.

Definition (Consensus Protocol)


(Informally), a consensus protocol is a system of n processes that communicate through a set of shared objects.

Propose

Communicate

Decide

More clarification on “termination”
termination requires wait-free implementation of the consensus object
binary consensus problem



Processes communicate through shared objects.

(redraw)

consensus object (fig here)

```
1  public interface Consensus<T> {  
2      T decide(T value);  
3  }
```


consensus protocol

```
1  public abstract class ConsensusProtocol<T>
2      implements Consensus<T> {
3      protected T[] proposed = (T[]) new Object[N];
4
5      void propose(T value) {
6          proposed[ThreadID.get()] = value;
7      }
8
9      public abstract T decide(T value);
10 }
```

implement X using Y

```
1 public class CASConsensus<T>
2     extends ConsensusProtocol<T> {
3     private final int FIRST = -1;
4     private AtomicInteger r = new AtomicInteger(FIRST);
5
6     @Override
7     public T decide(T value) {
8         propose(value);
9
10        int i = ThreadID.get();
11        if (r.compareAndSet(FIRST, i)) // I won
12            return proposed[i];
13        else // I lose
14            return proposed[r.get()];
15    }
16 }
```

Theorem (Computational Power of CAS)

*A register providing **compareAndSet()** and **get()** methods can solve the consensus problem for **any number** of threads.*

Theorem (Computational Power of CAS)

*A register providing **compareAndSet()** and **get()** methods can solve the consensus problem for **any number** of threads.*

In terms of “Consensus Number”:

Theorem (Consensus Number of CAS)

*A register providing **compareAndSet()** and **get()** methods has **infinite consensus number**.*

Definition (Solving Consensus)

“ X solves n -process consensus” if there exists a consensus protocol

$$P_1, \dots, P_n; W, X$$

- ▶ W is a set of (atomic) read/write registers;
- ▶ W and X may be initialized to any state.

Definition (Consensus Number)

The **consensus number** for X is the largest n for which X solves n -thread consensus.


If no largest n exists, the consensus number is said to be **infinite**.

Lemma (Y Implements X)

Theorem (Consensus Number as ...)

in the following, main results (table)

beautiful ideas and proofs



Protocol in terms of set:

Protocol

$$\mathcal{P} = \{\text{All executions of this protocol}\}$$

Execution

$$e = \sigma_0 \xrightarrow{o_1} \sigma_1 \xrightarrow{o_2} \cdots \xrightarrow{o_{n-1}} \sigma_n$$


State

σ_i : states of individual threads + states of shared objects

Operation

o_i : method calls to a shared object

Modeling \mathcal{P} as a Computation “Tree” (Binary Consensus for 2 Threads)



Theorem (Bivalent Initial State)

Every 2-thread binary consensus protocol has a bivalent initial state.

Theorem (Bivalent Initial State)


Every 2-thread binary consensus protocol has a bivalent initial state.

$$(A, B, O) : (0, 0, *) \quad (1, 1, *) \quad (0, 1, *) \quad (1, 0, *)$$

Theorem (Bivalent Initial State)

Every 2-thread binary consensus protocol has a bivalent initial state.

$$(A, B, O) : (0, 0, *) \quad (1, 1, *) \quad (0, 1, *) \quad (1, 0, *)$$



Lemma (Bivalent Initial State)

Every n -thread (binary) consensus protocol has a bivalent initial state.

Lemma (Bivalent Initial State)

Every n -thread (binary) consensus protocol has a bivalent initial state.

Theorem (Existence of Critical States)


Every wait-free consensus protocol has a critical state.

Lemma (Bivalent Initial State)

Every n -thread (binary) consensus protocol has a bivalent initial state.

Theorem (Existence of Critical States)

Every wait-free consensus protocol has a critical state.






Theorem (Consensus Number of Atomic Registers)


Atomic registers have consensus number 1.


Proof.

- (1) Run the protocol until it reaches a critical state s .
- (2) What is the next?









Theorem

Thank
You!