

Consensus Number

魏恒峰

hfwei@nju.edu.cn

2017 年 12 月 14 日




Feel Free to
Ask Questions



Are you Familiar with **Concurrent Programming**?



The Key: **Synchronization!**



Using the **Synchronization Primitives**
Provided by Your Favorite Languages.

synchronized
Semaphore

BlockingQueue
ConcurrentMap

Phaser
Barrier

synchronized
Semaphore

BlockingQueue
ConcurrentMap

Phaser
Barrier

```
class AtomicInteger:
    get()
    set(int newValue)

    getAndIncrement()
    getAndDecrement()
    getAndSet(int newValue)

    compareAndSet(int expectedValue, int newValue)
```

```
compareAndSet(int expectedValue, int newValue)  
compareAndSwap(int expectedValue, int newValue)
```

CAS — CMPXCHG

impl.

usage

Consensus



"It looks like we have a consensus."

"It looks like we have a consensus."

Propose



Decide

Propose



Decide


Definition (The Consensus Problem)

Agreement All (non-faulty) processes must **agree on the same value**.

Validity The common decision value must be the value **proposed** by some process.

Termination Each (non-faulty) process must **eventually** decide on a value.

More clarification on “termination”
binary consensus problem



(redraw)

consensus object (fig here)

```
1  public interface Consensus<T> {  
2      T decide(T value);  
3  }
```

consensus protocol

```
1  public abstract class ConsensusProtocol<T>
2      implements Consensus<T> {
3      protected T[] proposed = (T[]) new Object[N];
4
5      void propose(T value) {
6          proposed[ThreadID.get()] = value;
7      }
8
9      public abstract T decide(T value);
10 }
```

implement X using Y


```
1 public class CASConsensus<T>
2     extends ConsensusProtocol<T> {
3     private final int FIRST = -1;
4     private AtomicInteger r = new AtomicInteger(FIRST);
5
6     @Override
7     public T decide(T value) {
8         propose(value);
9
10        int i = ThreadID.get();
11        if (r.compareAndSet(FIRST, i)) // I won
12            return proposed[i];
13        else // I lose
14            return proposed[r.get()];
15    }
16 }
```

Theorem (Computational Power of CAS)

*A register providing **compareAndSet()** and **get()** methods can solve the consensus problem for **any number** of threads.*

Theorem (Computational Power of CAS)

*A register providing **compareAndSet()** and **get()** methods can solve the consensus problem for **any number** of threads.*

In terms of “Consensus Number”:

Theorem (Consensus Number of CAS)

*A register providing **compareAndSet()** and **get()** methods has **infinite consensus number**.*

Definition (Consensus Number)

The **consensus number** for X is the largest n for which X solves n -thread consensus.

If no largest n exists, the consensus number is said to be **infinite**.

Lemma (Y Implements X)

Theorem (Consensus Number as ...)

in the following, main results (table)

beautiful ideas and proofs

Thank
You!